

# SQL Interview Preparation Guide

## Complete Beginner to Advanced SQL Reference

This comprehensive guide covers everything you need to know about SQL for technical interviews:

- Database fundamentals and concepts
- All SQL commands from basic to advanced
- JOIN operations and subqueries
- Window functions and complex queries
- Performance optimization techniques
- 100+ practical examples and solutions
- Common interview questions and patterns

*Perfect for beginners who know nothing about SQL and want to master it for interviews!*

# Table of Contents

1. [SQL Basics](#sql-basics)
2. [Database Fundamentals](#database-fundamentals)
3. [Data Types](#data-types)
4. [SELECT Statement](#select-statement)
5. [Filtering Data](#filtering-data)
6. [Sorting and Grouping](#sorting-and-grouping)
7. [Joins](#joins)
8. [Subqueries](#subqueries)
9. [INSERT, UPDATE, DELETE](#insert-update-delete)
10. [CREATE and ALTER](#create-and-alter)
11. [Aggregate Functions](#aggregate-functions)
12. [Window Functions](#window-functions)
13. [String Functions](#string-functions)
14. [Date and Time Functions](#date-and-time-functions)
15. [Constraints](#constraints)
16. [Indexes](#indexes)
17. [Views](#views)
18. [Stored Procedures](#stored-procedures)
19. [Transactions](#transactions)
20. [Advanced SQL Concepts](#advanced-sql-concepts)
21. [Performance Optimization](#performance-optimization)
22. [Common Interview Questions](#common-interview-questions)

## 23. [Practice Problems](#practice-problems)

# SQL Basics

## What is SQL?

**SQL (Structured Query Language)** is a programming language designed for managing and manipulating relational databases. It's used to:

- Query data **FROM** databases
- **INSERT**, **UPDATE**, and **DELETE** records
- **CREATE** and modify **DATABASE** structures
- Control access to data

## Key SQL Commands Categories

1. **DDL (Data Definition Language)**: **CREATE**, **ALTER**, **DROP**
2. **DML (Data Manipulation Language)**: **SELECT**, **INSERT**, **UPDATE**, **DELETE**
3. **DCL (Data Control Language)**: GRANT, REVOKE
4. **TCL (Transaction Control Language)**: COMMIT, ROLLBACK, SAVEPOINT

## Database Fundamentals

*Understanding database concepts including tables, relationships, keys, and normalization principles.*

## What is a Database?

A **DATABASE** is an organized collection of structured data stored electronically in a computer system.

## What is a Table?

A **TABLE** is a collection of related data entries consisting of columns and rows.

## Key Terms

- **Row/Record:** A single entry in a **TABLE**
- **Column/Field:** A vertical entity in a **TABLE** containing all information associated with a specific field
- **Primary Key:** A unique identifier for each record in a **TABLE**
- **Foreign Key:** A field that links to the primary key of another **TABLE**
- **Schema:** The structure of a **DATABASE** (tables, columns, relationships)

## Sample Database Structure

```
-- Employees Table
CREATE TABLE employees (
    employee_id INT PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50),
    email VARCHAR(100),
    department_id INT,
    salary DECIMAL(10,2),
    hire_date DATE
);
```

```
-- Departments Table
CREATE TABLE departments (
    department_id INT PRIMARY KEY,
    department_name VARCHAR(100),
    manager_id INT
);
```

## Data Types

*SQL supports various data types including numeric, string, date/time, and boolean types.*

## Common SQL Data Types

### Numeric Types

INT	-- Integer numbers (-2,147,483,648 to 2,147,483,647)
BIGINT	-- Large integer numbers
DECIMAL(p,s)	-- Fixed-point numbers (p=precision, s=scale)
NUMERIC(p,s)	-- Same as DECIMAL
FLOAT	-- Floating-point numbers
REAL	-- Single precision floating-point

### String Types

CHAR(n)	-- Fixed-length string (n characters)
VARCHAR(n)	-- Variable-length string (up to n characters)
TEXT	-- Large variable-length string
NCHAR(n)	-- Fixed-length Unicode string
NVARCHAR(n)	-- Variable-length Unicode string

### Date and Time Types

DATE	-- Date (YYYY-MM-DD)
TIME	-- Time (HH:MM:SS)
DATETIME	-- Date and time (YYYY-MM-DD HH:MM:SS)
TIMESTAMP	-- Date and time with timezone

```
YEAR          -- Year (YYYY)
```

## Boolean Type

```
BOOLEAN      -- TRUE/FALSE values
BIT          -- 0 or 1
```

# SELECT Statement

*The most fundamental SQL command for retrieving data from database tables.*

## Basic SELECT Syntax

```
SELECT column1, column2, ...
FROM table_name;
```

## Examples

```
-- Select all columns
SELECT * FROM employees;

-- Select specific columns
SELECT first_name, last_name, salary FROM employees;

-- Select with alias
SELECT first_name AS "First Name",
       last_name AS "Last Name",
       salary AS "Annual Salary"
FROM employees;

-- Select distinct values
SELECT DISTINCT department_id FROM employees;

-- Select with calculations
SELECT first_name,
       last_name,
```

```
    salary,  
    salary * 12 AS annual_salary  
FROM employees;
```

## Filtering Data

*Using WHERE clauses and conditions to filter and retrieve specific data.*

## WHERE Clause

```
-- Basic WHERE  
SELECT * FROM employees WHERE department_id = 10;  
  
-- Multiple conditions with AND  
SELECT * FROM employees  
WHERE department_id = 10 AND salary > 50000;  
  
-- Multiple conditions with OR  
SELECT * FROM employees  
WHERE department_id = 10 OR department_id = 20;  
  
-- NOT operator  
SELECT * FROM employees  
WHERE NOT department_id = 10;
```

## Comparison Operators

```
-- Equal  
SELECT * FROM employees WHERE salary = 50000;  
  
-- Not equal  
SELECT * FROM employees WHERE salary != 50000;  
SELECT * FROM employees WHERE salary <> 50000;  
  
-- Greater than  
SELECT * FROM employees WHERE salary > 50000;  
  
-- Less than  
SELECT * FROM employees WHERE salary < 50000;
```

```
-- Greater than or equal  
SELECT * FROM employees WHERE salary >= 50000;  
  
-- Less than or equal  
SELECT * FROM employees WHERE salary <= 50000;
```

## BETWEEN Operator

```
-- Between (inclusive)  
SELECT * FROM employees WHERE salary BETWEEN 40000 AND 60000;  
  
-- Not between  
SELECT * FROM employees WHERE salary NOT BETWEEN 40000 AND 60000;
```

## IN Operator

```
-- IN operator  
SELECT * FROM employees WHERE department_id IN (10, 20, 30);  
  
-- NOT IN operator  
SELECT * FROM employees WHERE department_id NOT IN (10, 20, 30);
```

## LIKE Operator (Pattern Matching)

```
-- Starts with 'J'  
SELECT * FROM employees WHERE first_name LIKE 'J%';  
  
-- Ends with 'son'  
SELECT * FROM employees WHERE last_name LIKE '%son';  
  
-- Contains 'an'  
SELECT * FROM employees WHERE first_name LIKE '%an%';  
  
-- Second character is 'o'  
SELECT * FROM employees WHERE first_name LIKE '_o%';  
  
-- Exactly 5 characters  
SELECT * FROM employees WHERE first_name LIKE '_____';
```

## IS NULL and IS NOT NULL

```
-- Check for NULL values
SELECT * FROM employees WHERE manager_id IS NULL;

-- Check for NOT NULL values
SELECT * FROM employees WHERE manager_id IS NOT NULL;
```

## Sorting and Grouping

### ORDER BY Clause

```
-- Sort ascending (default)
SELECT * FROM employees ORDER BY salary;

-- Sort descending
SELECT * FROM employees ORDER BY salary DESC;

-- Sort by multiple columns
SELECT * FROM employees ORDER BY department_id, salary DESC;

-- Sort by column position
SELECT first_name, last_name, salary
FROM employees
ORDER BY 3 DESC; -- Sort by 3rd column (salary)
```

### GROUP BY Clause

```
-- Group by single column
SELECT department_id, COUNT(*) as employee_count
FROM employees
GROUP BY department_id;

-- Group by multiple columns
SELECT department_id,
       EXTRACT(YEAR FROM hire_date) as hire_year,
       COUNT(*) as employee_count
FROM employees
```

```
GROUP BY department_id, EXTRACT(YEAR FROM hire_date);
```

## HAVING Clause

```
-- HAVING (filter groups)
SELECT department_id, COUNT(*) as employee_count
FROM employees
GROUP BY department_id
HAVING COUNT(*) > 5;
```

```
-- HAVING with aggregate functions
SELECT department_id, AVG(salary) as avg_salary
FROM employees
GROUP BY department_id
HAVING AVG(salary) > 50000;
```

## Joins

### Types of Joins

1. **INNER JOIN**: Returns records that have matching values in both tables
2. **LEFT JOIN**: Returns all records **FROM** left **TABLE**, matched records **FROM** right **TABLE**
3. **RIGHT JOIN**: Returns all records **FROM** right **TABLE**, matched records **FROM** left **TABLE**
4. **FULL OUTER JOIN**: Returns all records when there's a match in either **TABLE**
5. **CROSS JOIN**: Returns Cartesian product of both tables

### INNER JOIN

```
-- Basic INNER JOIN
SELECT e.first_name, e.last_name, d.department_name
FROM employees e
```

```
INNER JOIN departments d ON e.department_id = d.department_id;

-- Multiple table join
SELECT e.first_name, e.last_name, d.department_name, p.project_name
FROM employees e
INNER JOIN departments d ON e.department_id = d.department_id
INNER JOIN projects p ON e.employee_id = p.employee_id;
```

## LEFT JOIN

```
-- LEFT JOIN (all employees, even without departments)
SELECT e.first_name, e.last_name, d.department_name
FROM employees e
LEFT JOIN departments d ON e.department_id = d.department_id;
```

## RIGHT JOIN

```
-- RIGHT JOIN (all departments, even without employees)
SELECT e.first_name, e.last_name, d.department_name
FROM employees e
RIGHT JOIN departments d ON e.department_id = d.department_id;
```

## FULL OUTER JOIN

```
-- FULL OUTER JOIN (all records from both tables)
SELECT e.first_name, e.last_name, d.department_name
FROM employees e
FULL OUTER JOIN departments d ON e.department_id = d.department_id;
```

## CROSS JOIN

```
-- CROSS JOIN (Cartesian product)
SELECT e.first_name, d.department_name
FROM employees e
CROSS JOIN departments d;
```

## Self Join

```
-- Self join (employees and their managers)
SELECT e.first_name AS employee, m.first_name AS manager
FROM employees e
LEFT JOIN employees m ON e.manager_id = m.employee_id;
```

## Subqueries

*Nested queries that provide powerful data retrieval and filtering capabilities.*

### What is a Subquery?

A **subquery** is a query nested inside another query. It can be used in **SELECT**, **INSERT**, **UPDATE**, or **DELETE** statements.

### Types of Subqueries

1. **Single-row subquery:** Returns one row
2. **Multi-row subquery:** Returns multiple rows
3. **Correlated subquery:** References columns **FROM** outer query

### Single-row Subquery

```
-- Find employees with salary higher than average
SELECT first_name, last_name, salary
FROM employees
WHERE salary > (SELECT AVG(salary) FROM employees);

-- Find employee with highest salary
SELECT first_name, last_name, salary
FROM employees
```

```
WHERE salary = (SELECT MAX(salary) FROM employees);
```

## Multi-row Subquery

```
-- Find employees in departments with more than 5 employees
SELECT first_name, last_name, department_id
FROM employees
WHERE department_id IN (
    SELECT department_id
    FROM employees
    GROUP BY department_id
    HAVING COUNT(*) > 5
);

-- Using ANY
SELECT first_name, last_name, salary
FROM employees
WHERE salary > ANY (
    SELECT salary
    FROM employees
    WHERE department_id = 10
);

-- Using ALL
SELECT first_name, last_name, salary
FROM employees
WHERE salary > ALL (
    SELECT salary
    FROM employees
    WHERE department_id = 10
);
```

## Correlated Subquery

```
-- Find employees earning more than average in their department
SELECT first_name, last_name, salary, department_id
FROM employees e1
WHERE salary > (
    SELECT AVG(salary)
    FROM employees e2
    WHERE e1.department_id = e2.department_id
);
```

## EXISTS and NOT EXISTS

```
-- Find employees who have at least one project
SELECT first_name, last_name
FROM employees e
WHERE EXISTS (
    SELECT 1
    FROM projects p
    WHERE p.employee_id = e.employee_id
);

-- Find employees who have no projects
SELECT first_name, last_name
FROM employees e
WHERE NOT EXISTS (
    SELECT 1
    FROM projects p
    WHERE p.employee_id = e.employee_id
);
```

## INSERT, UPDATE, DELETE

### INSERT Statement

```
-- Insert single row
INSERT INTO employees (employee_id, first_name, last_name, email, salary, hire_date)
VALUES (101, 'John', 'Doe', 'john.doe@company.com', 50000, '2023-01-15');

-- Insert multiple rows
INSERT INTO employees (employee_id, first_name, last_name, email, salary, hire_date)
VALUES
    (102, 'Jane', 'Smith', 'jane.smith@company.com', 55000, '2023-02-01'),
    (103, 'Bob', 'Johnson', 'bob.johnson@company.com', 52000, '2023-02-15');

-- Insert from another table
INSERT INTO employees_backup
SELECT * FROM employees WHERE department_id = 10;

-- Insert with subquery
INSERT INTO high_earners (employee_id, first_name, last_name, salary)
SELECT employee_id, first_name, last_name, salary
FROM employees
```

```
WHERE salary > (SELECT AVG(salary) FROM employees);
```

## UPDATE Statement

```
-- Update single column
UPDATE employees
SET salary = 55000
WHERE employee_id = 101;

-- Update multiple columns
UPDATE employees
SET salary = 60000, department_id = 20
WHERE employee_id = 101;

-- Update with subquery
UPDATE employees
SET salary = salary * 1.1
WHERE department_id = (
    SELECT department_id
    FROM departments
    WHERE department_name = 'Sales'
);

-- Update with JOIN
UPDATE employees e
SET salary = salary * 1.05
FROM departments d
WHERE e.department_id = d.department_id
AND d.department_name = 'IT';
```

## DELETE Statement

```
-- Delete specific records
DELETE FROM employees
WHERE employee_id = 101;

-- Delete with condition
DELETE FROM employees
WHERE department_id = 30 AND salary < 40000;

-- Delete with subquery
DELETE FROM employees
WHERE department_id IN (
    SELECT department_id
    FROM departments
```

```

        WHERE department_name = 'Temp'
);

-- Delete all records (keep table structure)
DELETE FROM employees;

```

## CREATE and ALTER

### CREATE TABLE

```

-- Basic CREATE TABLE
CREATE TABLE employees (
    employee_id INT PRIMARY KEY,
    first_name VARCHAR(50) NOT NULL,
    last_name VARCHAR(50) NOT NULL,
    email VARCHAR(100) UNIQUE,
    department_id INT,
    salary DECIMAL(10,2),
    hire_date DATE,
    FOREIGN KEY (department_id) REFERENCES departments(department_id)
);

```

```

-- CREATE TABLE with constraints
CREATE TABLE products (
    product_id INT AUTO_INCREMENT PRIMARY KEY,
    product_name VARCHAR(100) NOT NULL,
    price DECIMAL(8,2) CHECK (price > 0),
    category_id INT,
    created_date TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    INDEX idx_category (category_id)
);

```

```

-- CREATE TABLE from another table
CREATE TABLE employees_backup AS
SELECT * FROM employees;

```

```

-- CREATE TABLE with specific columns
CREATE TABLE employee_summary AS
SELECT employee_id, first_name, last_name, salary
FROM employees
WHERE department_id = 10;

```

## ALTER TABLE

```
-- Add column
ALTER TABLE employees
ADD COLUMN phone VARCHAR(20);

-- Drop column
ALTER TABLE employees
DROP COLUMN phone;

-- Modify column
ALTER TABLE employees
MODIFY COLUMN salary DECIMAL(12,2);

-- Rename column
ALTER TABLE employees
RENAME COLUMN first_name TO fname;

-- Add constraint
ALTER TABLE employees
ADD CONSTRAINT fk_department
FOREIGN KEY (department_id) REFERENCES departments(department_id);

-- Drop constraint
ALTER TABLE employees
DROP CONSTRAINT fk_department;

-- Rename table
ALTER TABLE employees
RENAME TO staff;
```

## DROP TABLE

```
-- Drop table
DROP TABLE employees_backup;

-- Drop table if exists
DROP TABLE IF EXISTS employees_backup;
```

## Aggregate Functions

*Built-in functions like COUNT, SUM, AVG for data analysis and calculations.*

## Common Aggregate Functions

```
-- COUNT: Count number of rows
SELECT COUNT(*) FROM employees;
SELECT COUNT(DISTINCT department_id) FROM employees;

-- SUM: Sum of values
SELECT SUM(salary) FROM employees;
SELECT SUM(salary) FROM employees WHERE department_id = 10;

-- AVG: Average of values
SELECT AVG(salary) FROM employees;
SELECT AVG(salary) FROM employees WHERE department_id = 10;

-- MIN: Minimum value
SELECT MIN(salary) FROM employees;
SELECT MIN(hire_date) FROM employees;

-- MAX: Maximum value
SELECT MAX(salary) FROM employees;
SELECT MAX(hire_date) FROM employees;
```

## Aggregate Functions with GROUP BY

```
-- Count employees by department
SELECT department_id, COUNT(*) as employee_count
FROM employees
GROUP BY department_id;

-- Average salary by department
SELECT department_id, AVG(salary) as avg_salary
FROM employees
GROUP BY department_id;

-- Multiple aggregates
SELECT department_id,
       COUNT(*) as employee_count,
       AVG(salary) as avg_salary,
       MIN(salary) as min_salary,
       MAX(salary) as max_salary,
       SUM(salary) as total_salary
FROM employees
GROUP BY department_id;
```

# Window Functions

*Advanced functions for performing calculations across related rows.*

## What are Window Functions?

**Window functions** perform calculations across a set of rows related to the current row, without collapsing the rows into a single result.

### ROW\_NUMBER()

```
-- Add row numbers
SELECT employee_id, first_name, last_name, salary,
       ROW_NUMBER() OVER (ORDER BY salary DESC) as row_num
FROM employees;

-- Row numbers within each department
SELECT employee_id, first_name, last_name, department_id, salary,
       ROW_NUMBER() OVER (PARTITION BY department_id ORDER BY salary DESC) as dep
FROM employees;
```

### RANK() and DENSE\_RANK()

```
-- RANK (with gaps)
SELECT employee_id, first_name, last_name, salary,
       RANK() OVER (ORDER BY salary DESC) as salary_rank
FROM employees;

-- DENSE_RANK (no gaps)
SELECT employee_id, first_name, last_name, salary,
       DENSE_RANK() OVER (ORDER BY salary DESC) as salary_dense_rank
FROM employees;
```

## LAG() and LEAD()

```
-- LAG (previous row value)
SELECT employee_id, first_name, salary,
       LAG(salary, 1) OVER (ORDER BY salary) as prev_salary
FROM employees;

-- LEAD (next row value)
SELECT employee_id, first_name, salary,
       LEAD(salary, 1) OVER (ORDER BY salary) as next_salary
FROM employees;
```

## FIRST\_VALUE() and LAST\_VALUE()

```
-- FIRST_VALUE and LAST_VALUE
SELECT employee_id, first_name, salary,
       FIRST_VALUE(salary) OVER (ORDER BY salary DESC) as highest_salary,
       LAST_VALUE(salary) OVER (ORDER BY salary DESC
                                ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING)
FROM employees;
```

## SUM() and AVG() as Window Functions

```
-- Running total
SELECT employee_id, first_name, salary,
       SUM(salary) OVER (ORDER BY employee_id) as running_total
FROM employees;

-- Moving average
SELECT employee_id, first_name, salary,
       AVG(salary) OVER (ORDER BY employee_id ROWS BETWEEN 2 PRECEDING AND CURRENT ROW)
FROM employees;
```

## String Functions

*Functions for manipulating and processing text data in SQL.*

## Common String Functions

```
-- UPPER: Convert to uppercase
SELECT UPPER(first_name) FROM employees;

-- LOWER: Convert to lowercase
SELECT LOWER(first_name) FROM employees;

-- LENGTH: Get string length
SELECT first_name, LENGTH(first_name) as name_length FROM employees;

-- SUBSTRING: Extract substring
SELECT SUBSTRING(first_name, 1, 3) as first_three FROM employees;

-- CONCAT: Concatenate strings
SELECT CONCAT(first_name, ' ', last_name) as full_name FROM employees;

-- TRIM: Remove leading/trailing spaces
SELECT TRIM(first_name) FROM employees;

-- LTRIM and RTRIM: Remove left/right spaces
SELECT LTRIM(first_name), RTRIM(last_name) FROM employees;

-- REPLACE: Replace substring
SELECT REPLACE(email, '@company.com', '@newcompany.com') FROM employees;

-- LEFT and RIGHT: Extract from left/right
SELECT LEFT(first_name, 3), RIGHT(last_name, 3) FROM employees;

-- CHARINDEX/POSITION: Find position of substring
SELECT CHARINDEX('@', email) as at_position FROM employees;
```

## String Functions Examples

```
-- Extract domain from email
SELECT email,
       SUBSTRING(email, CHARINDEX('@', email) + 1, LEN(email)) as domain
FROM employees;

-- Create initials
SELECT CONCAT(LEFT(first_name, 1), '.', LEFT(last_name, 1), '.') as initials
FROM employees;

-- Check if email is valid
SELECT email,
       CASE WHEN email LIKE '%@%.%' THEN 'Valid' ELSE 'Invalid' END as email_stat
FROM employees;
```

# Date and Time Functions

*Functions for working with date and time data types.*

## Current Date/Time Functions

```
-- Current date and time
SELECT GETDATE();          -- SQL Server
SELECT NOW();              -- MySQL
SELECT SYSDATE;            -- Oracle
SELECT CURRENT_TIMESTAMP;  -- Standard SQL

-- Current date only
SELECT CAST(GETDATE() AS DATE);  -- SQL Server
SELECT CURDATE();              -- MySQL
SELECT CURRENT_DATE;           -- Standard SQL

-- Current time only
SELECT CAST(GETDATE() AS TIME);  -- SQL Server
SELECT CURTIME();              -- MySQL
SELECT CURRENT_TIME;           -- Standard SQL
```

## Date Arithmetic

```
-- Add days
SELECT DATEADD(DAY, 30, hire_date) as thirty_days_later FROM employees;

-- Subtract days
SELECT DATEADD(DAY, -30, hire_date) as thirty_days_earlier FROM employees;

-- Add months
SELECT DATEADD(MONTH, 6, hire_date) as six_months_later FROM employees;

-- Add years
SELECT DATEADD(YEAR, 1, hire_date) as one_year_later FROM employees;
```

## Date Difference

```
-- Difference in days
SELECT first_name, hire_date,
       DATEDIFF(DAY, hire_date, GETDATE()) as days_employed
FROM employees;

-- Difference in months
SELECT first_name, hire_date,
       DATEDIFF(MONTH, hire_date, GETDATE()) as months_employed
FROM employees;

-- Difference in years
SELECT first_name, hire_date,
       DATEDIFF(YEAR, hire_date, GETDATE()) as years_employed
FROM employees;
```

## Extract Date Parts

```
-- Extract year, month, day
SELECT hire_date,
       YEAR(hire_date) as hire_year,
       MONTH(hire_date) as hire_month,
       DAY(hire_date) as hire_day
FROM employees;

-- Extract weekday
SELECT hire_date,
       DATENAME(WEEKDAY, hire_date) as hire_weekday
FROM employees;

-- Extract quarter
SELECT hire_date,
       DATEPART(QUARTER, hire_date) as hire_quarter
FROM employees;
```

## Constraints

*Rules and restrictions to ensure data integrity and consistency.*

# Types of Constraints

1. **NOT NULL**: Ensures column cannot have NULL values
2. **UNIQUE**: Ensures all values in column are unique
3. **PRIMARY KEY**: Combination of NOT NULL and UNIQUE
4. **FOREIGN KEY**: Links to primary key of another **TABLE**
5. **CHECK**: Ensures values meet specific condition
6. **DEFAULT**: Sets default value for column

# Constraint Examples

```
-- Create table with constraints
CREATE TABLE employees (
    employee_id INT PRIMARY KEY,
    first_name VARCHAR(50) NOT NULL,
    last_name VARCHAR(50) NOT NULL,
    email VARCHAR(100) UNIQUE,
    department_id INT,
    salary DECIMAL(10,2) CHECK (salary > 0),
    hire_date DATE DEFAULT CURRENT_DATE,
    status VARCHAR(20) DEFAULT 'Active' CHECK (status IN ('Active', 'Inactive')),
    FOREIGN KEY (department_id) REFERENCES departments(department_id)
);

-- Add constraints to existing table
ALTER TABLE employees
ADD CONSTRAINT chk_salary CHECK (salary >= 0);

ALTER TABLE employees
ADD CONSTRAINT fk_department
FOREIGN KEY (department_id) REFERENCES departments(department_id);

-- Drop constraints
ALTER TABLE employees
DROP CONSTRAINT chk_salary;
```

# Indexes

*Database structures that improve query performance and data retrieval speed.*

## What are Indexes?

**Indexes** are **DATABASE** objects that improve the speed of data retrieval operations on a **TABLE**.

## Types of Indexes

1. **Clustered INDEX:** Determines physical order of data
2. **Non-Clustered INDEX:** Separate structure pointing to data
3. **Unique INDEX:** Ensures uniqueness and improves performance
4. **Composite INDEX:** **INDEX** on multiple columns

## Creating Indexes

```
-- Create simple index
CREATE INDEX idx_employee_lastname ON employees(last_name);

-- Create composite index
CREATE INDEX idx_employee_dept_salary ON employees(department_id, salary);

-- Create unique index
CREATE UNIQUE INDEX idx_employee_email ON employees(email);

-- Create clustered index
CREATE CLUSTERED INDEX idx_employee_id ON employees(employee_id);
```

## Managing Indexes

```
-- View indexes
```

```

SELECT * FROM sys.indexes WHERE object_id = OBJECT_ID('employees');

-- Drop index
DROP INDEX idx_employee_lastname ON employees;

-- Rebuild index
ALTER INDEX idx_employee_lastname ON employees REBUILD;

```

## Views

*Virtual tables that provide customized data presentation and security.*

### What are Views?

A **VIEW** is a virtual **TABLE** based on the result of a SQL statement. It contains rows and columns like a real **TABLE**.

### Creating Views

```

-- Simple view
CREATE VIEW employee_details AS
SELECT e.employee_id, e.first_name, e.last_name, d.department_name, e.salary
FROM employees e
JOIN departments d ON e.department_id = d.department_id;

-- View with calculations
CREATE VIEW employee_summary AS
SELECT department_id,
       COUNT(*) as employee_count,
       AVG(salary) as avg_salary,
       MAX(salary) as max_salary,
       MIN(salary) as min_salary
FROM employees
GROUP BY department_id;

-- View with filtering
CREATE VIEW high_earners AS
SELECT employee_id, first_name, last_name, salary
FROM employees
WHERE salary > 75000;

```

## Using Views

```
-- Query from view
SELECT * FROM employee_details;

-- Filter view results
SELECT * FROM employee_details WHERE department_name = 'Sales';

-- Join views
SELECT ed.first_name, ed.last_name, es.avg_salary
FROM employee_details ed
JOIN employee_summary es ON ed.department_id = es.department_id;
```

## Managing Views

```
-- Modify view
ALTER VIEW employee_details AS
SELECT e.employee_id, e.first_name, e.last_name, d.department_name, e.salary, e.h
FROM employees e
JOIN departments d ON e.department_id = d.department_id;

-- Drop view
DROP VIEW employee_details;
```

## Stored Procedures

*Reusable database programs that encapsulate complex operations.*

## What are Stored Procedures?

A **stored PROCEDURE** is a prepared SQL code that can be saved and reused.

## Creating Stored Procedures

```
-- Simple stored procedure
CREATE PROCEDURE GetEmployeeCount
AS
BEGIN
    SELECT COUNT(*) as total_employees FROM employees;
END;

-- Stored procedure with parameters
CREATE PROCEDURE GetEmployeesByDepartment
    @DepartmentId INT
AS
BEGIN
    SELECT employee_id, first_name, last_name, salary
    FROM employees
    WHERE department_id = @DepartmentId;
END;

-- Stored procedure with output parameter
CREATE PROCEDURE GetEmployeeCountByDept
    @DepartmentId INT,
    @EmployeeCount INT OUTPUT
AS
BEGIN
    SELECT @EmployeeCount = COUNT(*)
    FROM employees
    WHERE department_id = @DepartmentId;
END;
```

## Executing Stored Procedures

```
-- Execute simple procedure
EXEC GetEmployeeCount;

-- Execute with parameters
EXEC GetEmployeesByDepartment @DepartmentId = 10;

-- Execute with output parameter
DECLARE @Count INT;
EXEC GetEmployeeCountByDept @DepartmentId = 10, @EmployeeCount = @Count OUTPUT;
SELECT @Count as EmployeeCount;
```

# Transactions

*Mechanisms for ensuring data consistency and handling concurrent operations.*

## What are Transactions?

A **transaction** is a sequence of SQL operations treated as a single unit of work.

## ACID Properties

1. **Atomicity:** All operations succeed or all fail
2. **Consistency:** **DATABASE** remains in valid state
3. **Isolation:** Concurrent transactions don't interfere
4. **Durability:** Committed changes persist

## Transaction Control

```
-- Basic transaction
BEGIN TRANSACTION;
    INSERT INTO employees VALUES (104, 'Alice', 'Wilson', 'alice@company.com', 10);
    UPDATE employees SET salary = salary * 1.1 WHERE department_id = 10;
COMMIT;

-- Transaction with rollback
BEGIN TRANSACTION;
    DELETE FROM employees WHERE employee_id = 104;
    -- Check if we accidentally deleted too many records
    IF @@ROWCOUNT > 1
        BEGIN
            ROLLBACK;
            PRINT 'Transaction rolled back - too many records affected';
        END
    ELSE
        BEGIN
            COMMIT;
            PRINT 'Transaction committed successfully';
        END;
END;
```

```
-- Using savepoints
BEGIN TRANSACTION;
    INSERT INTO employees VALUES (105, 'Bob', 'Brown', 'bob@company.com', 20, 55000);
    SAVE TRANSACTION SavePoint1;

    UPDATE employees SET salary = 70000 WHERE employee_id = 105;

    -- Something went wrong, rollback to savepoint
    ROLLBACK TRANSACTION SavePoint1;

    -- Continue with original transaction
    UPDATE employees SET salary = 58000 WHERE employee_id = 105;
COMMIT;
```

## Advanced SQL Concepts

*Complex topics including CTEs, recursive queries, and advanced optimization.*

## Common Table Expressions (CTEs)

```
-- Simple CTE
WITH DepartmentStats AS (
    SELECT department_id,
           COUNT(*) as employee_count,
           AVG(salary) as avg_salary
    FROM employees
   GROUP BY department_id
)
SELECT d.department_name, ds.employee_count, ds.avg_salary
  FROM DepartmentStats ds
 JOIN departments d ON ds.department_id = d.department_id;

-- Recursive CTE (Employee hierarchy)
WITH EmployeeHierarchy AS (
    -- Anchor member (top-level managers)
    SELECT employee_id, first_name, last_name, manager_id, 0 as level
      FROM employees
     WHERE manager_id IS NULL

    UNION ALL

    -- Recursive member
    SELECT e.employee_id, e.first_name, e.last_name, e.manager_id, eh.level + 1
      FROM employees e
      JOIN EmployeeHierarchy eh
        ON e.manager_id = eh.employee_id
)
```

```

        FROM employees e
        JOIN EmployeeHierarchy eh ON e.manager_id = eh.employee_id
    )
SELECT * FROM EmployeeHierarchy ORDER BY level, employee_id;

```

## PIVOT and UNPIVOT

```

-- PIVOT example
SELECT *
FROM (
    SELECT department_name, YEAR(hire_date) as hire_year
    FROM employees e
    JOIN departments d ON e.department_id = d.department_id
) as SourceTable
PIVOT (
    COUNT(hire_year)
    FOR hire_year IN ([2020], [2021], [2022], [2023])
) as PivotTable;

-- UNPIVOT example
SELECT department_name, hire_year, employee_count
FROM (
    SELECT department_name, [2020], [2021], [2022], [2023]
    FROM PivotTable
) as SourceTable
UNPIVOT (
    employee_count FOR hire_year IN ([2020], [2021], [2022], [2023])
) as UnpivotTable;

```

## CASE Statements

```

-- Simple CASE
SELECT first_name, last_name, salary,
CASE
    WHEN salary < 40000 THEN 'Low'
    WHEN salary BETWEEN 40000 AND 70000 THEN 'Medium'
    ELSE 'High'
END as salary_grade
FROM employees;

-- CASE with multiple conditions
SELECT first_name, last_name, department_id, salary,
CASE
    WHEN department_id = 10 AND salary > 60000 THEN 'Senior IT'
    WHEN department_id = 10 THEN 'IT'

```

```
        WHEN department_id = 20 AND salary > 50000 THEN 'Senior Sales'
        WHEN department_id = 20 THEN 'Sales'
        ELSE 'Other'
    END as employee_category
FROM employees;
```

## Performance Optimization

*Techniques for improving query performance and database efficiency.*

### Query Optimization Tips

1. Use indexes appropriately
2. Avoid **SELECT** `\*` - specify needed columns
3. Use **WHERE** clauses to filter early
4. Use **JOINS** instead of subqueries when possible
5. Avoid functions in **WHERE** clauses
6. Use **LIMIT/TOP** for large datasets

### Execution Plan Analysis

```
-- View execution plan (SQL Server)
SET SHOWPLAN_ALL ON;
SELECT * FROM employees WHERE department_id = 10;
SET SHOWPLAN_ALL OFF;

-- Include actual execution plan
SET STATISTICS IO ON;
SET STATISTICS TIME ON;
SELECT * FROM employees WHERE department_id = 10;
SET STATISTICS IO OFF;
SET STATISTICS TIME OFF;
```

## Optimization Examples

```
-- Bad: Function in WHERE clause
SELECT * FROM employees WHERE YEAR(hire_date) = 2023;

-- Good: Range condition
SELECT * FROM employees WHERE hire_date >= '2023-01-01' AND hire_date < '2024-01-01';

-- Bad: SELECT *
SELECT * FROM employees WHERE department_id = 10;

-- Good: Specific columns
SELECT employee_id, first_name, last_name, salary
FROM employees WHERE department_id = 10;

-- Bad: Correlated subquery
SELECT e1.first_name, e1.last_name
FROM employees e1
WHERE e1.salary > (SELECT AVG(e2.salary) FROM employees e2 WHERE e2.department_id = 10);

-- Good: Window function
SELECT first_name, last_name
FROM (
    SELECT first_name, last_name, salary,
           AVG(salary) OVER (PARTITION BY department_id) as dept_avg
    FROM employees
) t
WHERE salary > dept_avg;
```

## Common Interview Questions

*Frequently asked SQL questions and their solutions for technical interviews.*

### 1. Find Second Highest Salary

```
-- Method 1: Using window function
SELECT salary
FROM (
    SELECT salary, ROW_NUMBER() OVER (ORDER BY salary DESC) as rn
```

```

        FROM employees
) t
WHERE rn = 2;

-- Method 2: Using subquery
SELECT MAX(salary) as second_highest
FROM employees
WHERE salary < (SELECT MAX(salary) FROM employees);

-- Method 3: Using LIMIT/TOP
SELECT DISTINCT salary
FROM employees
ORDER BY salary DESC
LIMIT 1 OFFSET 1; -- MySQL
-- OR
SELECT TOP 1 salary
FROM (
    SELECT DISTINCT TOP 2 salary
    FROM employees
    ORDER BY salary DESC
) t
ORDER BY salary ASC; -- SQL Server

```

## 2. Find Nth Highest Salary

```

-- Using window function
SELECT salary
FROM (
    SELECT salary, DENSE_RANK() OVER (ORDER BY salary DESC) as rank
    FROM employees
) t
WHERE rank = N;

-- Using subquery (for 3rd highest)
SELECT salary
FROM employees e1
WHERE 3 = (
    SELECT COUNT(DISTINCT salary)
    FROM employees e2
    WHERE e2.salary >= e1.salary
);

```

## 3. Find Employees with Same Salary

```
SELECT e1.first_name, e1.last_name, e1.salary
```

```

FROM employees e1
JOIN employees e2 ON e1.salary = e2.salary AND e1.employee_id != e2.employee_id;

-- Or using window function
SELECT first_name, last_name, salary
FROM (
    SELECT first_name, last_name, salary,
           COUNT(*) OVER (PARTITION BY salary) as salary_count
    FROM employees
) t
WHERE salary_count > 1;

```

## 4. Find Employees with No Manager

```

SELECT first_name, last_name
FROM employees
WHERE manager_id IS NULL;

```

## 5. Find Department with Highest Average Salary

```

SELECT d.department_name, AVG(e.salary) as avg_salary
FROM employees e
JOIN departments d ON e.department_id = d.department_id
GROUP BY d.department_id, d.department_name
ORDER BY avg_salary DESC
LIMIT 1;

```

## 6. Find Employees Hired in Last 30 Days

```

SELECT first_name, last_name, hire_date
FROM employees
WHERE hire_date >= DATEADD(DAY, -30, GETDATE()); -- SQL Server
-- OR
WHERE hire_date >= DATE_SUB(CURDATE(), INTERVAL 30 DAY); -- MySQL

```

## 7. Find Duplicate Records

```
-- Find duplicate emails
```

```

SELECT email, COUNT(*) as count
FROM employees
GROUP BY email
HAVING COUNT(*) > 1;

-- Find all duplicate records
SELECT *
FROM employees
WHERE email IN (
    SELECT email
    FROM employees
    GROUP BY email
    HAVING COUNT(*) > 1
);

```

## 8. Delete Duplicate Records

```

-- Using window function
WITH DuplicateRecords AS (
    SELECT *,
        ROW_NUMBER() OVER (PARTITION BY email ORDER BY employee_id) as rn
    FROM employees
)
DELETE FROM DuplicateRecords
WHERE rn > 1;

```

## 9. Find Running Total

```

SELECT employee_id, first_name, salary,
       SUM(salary) OVER (ORDER BY employee_id) as running_total
FROM employees;

```

## 10. Find Top 3 Employees by Salary in Each Department

```

SELECT department_id, first_name, last_name, salary
FROM (
    SELECT department_id, first_name, last_name, salary,
           ROW_NUMBER() OVER (PARTITION BY department_id ORDER BY salary DESC) as rn
    FROM employees
) t
WHERE rn <= 3;

```

# Practice Problems

## Easy Level

1. List all employees with their department names
2. Find total number of employees in each department
3. Find employees with salary greater than 50000
4. List employees hired in 2023
5. Find the average salary of all employees

## Medium Level

1. Find employees earning more than their department average
2. List departments with more than 5 employees
3. Find employees who have been with company for more than 2 years
4. Calculate the percentage of total salary each employee represents
5. Find employees with same first name

## Hard Level

1. Find the 3rd highest salary without using TOP/LIMIT
2. **CREATE** a hierarchical query showing employee-manager relationships
3. Find departments **WHERE** average salary is above company average

**4. Calculate running totals and moving averages**

**5. Find employees whose salary is above 80th percentile**

## Solutions to Practice Problems

### Easy Level Solutions

```
-- 1. List all employees with their department names
SELECT e.first_name, e.last_name, d.department_name
FROM employees e
JOIN departments d ON e.department_id = d.department_id;

-- 2. Find total number of employees in each department
SELECT d.department_name, COUNT(*) as employee_count
FROM employees e
JOIN departments d ON e.department_id = d.department_id
GROUP BY d.department_id, d.department_name;

-- 3. Find employees with salary greater than 50000
SELECT first_name, last_name, salary
FROM employees
WHERE salary > 50000;

-- 4. List employees hired in 2023
SELECT first_name, last_name, hire_date
FROM employees
WHERE YEAR(hire_date) = 2023;

-- 5. Find the average salary of all employees
SELECT AVG(salary) as average_salary
FROM employees;
```

### Medium Level Solutions

```
-- 1. Find employees earning more than their department average
SELECT e.first_name, e.last_name, e.salary, e.department_id
FROM employees e
WHERE e.salary > (
    SELECT AVG(e2.salary)
    FROM employees e2
    WHERE e2.department_id = e.department_id
);
```

```

-- 2. List departments with more than 5 employees
SELECT d.department_name, COUNT(*) as employee_count
FROM employees e
JOIN departments d ON e.department_id = d.department_id
GROUP BY d.department_id, d.department_name
HAVING COUNT(*) > 5;

-- 3. Find employees who have been with company for more than 2 years
SELECT first_name, last_name, hire_date,
       DATEDIFF(YEAR, hire_date, GETDATE()) as years_employed
FROM employees
WHERE DATEDIFF(YEAR, hire_date, GETDATE()) > 2;

-- 4. Calculate the percentage of total salary each employee represents
SELECT first_name, last_name, salary,
       (salary * 100.0 / (SELECT SUM(salary) FROM employees)) as salary_percentag
FROM employees;

-- 5. Find employees with same first name
SELECT first_name, COUNT(*) as count
FROM employees
GROUP BY first_name
HAVING COUNT(*) > 1;

```

## SQL Best Practices

### 1. Naming Conventions

- Use lowercase with underscores for **TABLE** and column names
- Use descriptive names
- Prefix foreign keys with **TABLE** name
- Use singular nouns for **TABLE** names

### 2. Query Writing

- Use proper indentation and formatting

- Use **TABLE** aliases for readability
- Always specify column names in **INSERT** statements
- Use EXISTS instead of IN for better performance

### 3. Performance Tips

- **CREATE** indexes on frequently queried columns
- Use appropriate data types
- Avoid **SELECT \*** in production code
- Use LIMIT/TOP for large result sets
- Use connection pooling

### 4. Security

- Use parameterized queries to prevent SQL injection
- Grant minimum necessary permissions
- Use stored procedures for complex operations
- Regularly backup your **DATABASE**

## Conclusion

This SQL guide covers everything **FROM** basic concepts to advanced techniques needed for technical interviews. Practice these concepts regularly and work on real-world problems to become proficient in SQL.

## Key Takeaways

1. **Master the basics:** `SELECT, INSERT, UPDATE, DELETE, JOIN`
2. **Understand data relationships:** Primary keys, foreign keys, normalization
3. **Practice complex queries:** Subqueries, window functions, CTEs
4. **Learn optimization:** Indexes, execution plans, performance tuning
5. **Solve problems:** Work through common interview questions

## Next Steps

1. Practice on online platforms (LeetCode, HackerRank, SQLBolt)
2. Work with real databases
3. Learn **DATABASE**-specific features (MySQL, PostgreSQL, SQL Server)
4. Study **DATABASE** design and normalization
5. Explore advanced topics (triggers, functions, data warehousing)

Good luck with your SQL interviews! ■